

Table of Contents

- [Introduction to Pandas](#)
 - [Pandas Data Structures](#)
 - [Series](#)
 - [DataFrame](#)
 - [Exercise 1](#)
 - [Exercise 2](#)
 - [Importing data](#)
 - [Missing Values](#)
 - [Microsoft Excel](#)
 - [Indexing and Selection](#)
 - [Exercise 3](#)
 - [Exercise 4](#)
 - [Hierarchical indexing](#)
 - [Operations](#)
 - [Sorting](#)
 - [Exercise 5](#)
 - [Missing data](#)
 - [Data summarization](#)
 - [Plotting](#)
 - [Writing Data to Files](#)
 - [References](#)

Introduction to Pandas

pandas is a Python package providing fast, flexible, and expressive data structures designed to work with *relational* or *labeled* data both. It is a fundamental high-level building block for doing practical, real world data analysis in Python.

pandas is well suited for:

- Tabular data with heterogeneously-typed columns, as in an SQL table or Excel spreadsheet
- Ordered and unordered (not necessarily fixed-frequency) time series data.
- Arbitrary matrix data (homogeneously typed or heterogeneous) with row and column labels
- Any other form of observational / statistical data sets. The data actually need not be labeled at all to be placed into a pandas data structure

Key features:

- Easy handling of **missing data**
- **Size mutability**: columns can be inserted and deleted from DataFrame and higher dimensional objects
- Automatic and explicit **data alignment**: objects can be explicitly aligned to a set of labels, or the data can be aligned automatically

- Powerful, flexible **group by functionality** to perform split-apply-combine operations on data sets
- Intelligent label-based **slicing, fancy indexing, and subsetting** of large data sets
- Intuitive **merging and joining** data sets
- Flexible **reshaping and pivoting** of data sets
- **Hierarchical labeling** of axes
- Robust **IO tools** for loading data from flat files, Excel files, databases, and HDF5
- **Time series functionality**: date range generation and frequency conversion, moving window statistics, moving window linear regressions, date shifting and lagging, etc.

```
In [1]: import pandas as pd
import numpy as np
pd.options.mode.chained_assignment = None # default='warn', Mutes warnings when copying
```

Pandas Data Structures

Series

A **Series** is a single vector of data (like a NumPy 1-d array) with an *index* that labels each element in the vector.

```
In [2]: counts = pd.Series([632, 1638, 569, 115])
counts
```

```
Out[2]: 0      632
1     1638
2      569
3      115
dtype: int64
```

If an **index** is not specified, a default sequence of integers is assigned as the index. A NumPy array comprises the values of the **Series**, while the index is a pandas **Index** object.

```
In [3]: counts.values
```

```
Out[3]: array([ 632, 1638,  569,  115])
```

```
In [4]: counts.index
```

```
Out[4]: RangeIndex(start=0, stop=4, step=1)
```

We can assign meaningful labels to the index, if they are available:

```
In [5]: bacteria = pd.Series([632, 1638, 569, 115],
    index=['Firmicutes', 'Proteobacteria', 'Actinobacteria', 'Bacteroidetes'])

bacteria
```

```
Out[5]: Firmicutes      632
Proteobacteria    1638
Actinobacteria     569
Bacteroidetes     115
dtype: int64
```

These labels can be used to refer to the values in the **Series**.

```
In [6]: bacteria['Actinobacteria'] #dictionary style
```

Out[6]: 569

```
In [7]: bacteria[[name.endswith('bacteria') for name in bacteria.index]]
```

Out[7]:

| | |
|----------------|------|
| Proteobacteria | 1638 |
| Actinobacteria | 569 |

dtype: int64

```
In [8]: [name.endswith('bacteria') for name in bacteria.index] #mask
```

Out[8]: [False, True, True, False]

Notice that the indexing operation preserved the association between the values and the corresponding indices.

We can still use **positional indexing** if we wish.

```
In [9]: bacteria[0]
```

Out[9]: 632

We can give both the array of values and the index **meaningful labels** themselves:

```
In [10]: bacteria.name = 'counts'
bacteria.index.name = 'phylum'
bacteria
```

Out[10]:

| phylum | |
|----------------|------|
| Firmicutes | 632 |
| Proteobacteria | 1638 |
| Actinobacteria | 569 |
| Bacteroidetes | 115 |

Name: counts, dtype: int64

NumPy's math functions and other operations can be applied to Series without losing the data structure.

```
In [10]: np.log(bacteria)
```

Out[10]:

| | |
|----------------|----------|
| Firmicutes | 6.448889 |
| Proteobacteria | 7.401231 |
| Actinobacteria | 6.343880 |
| Bacteroidetes | 4.744932 |

dtype: float64

```
In [11]: bacteria.apply(np.log)
```

Out[11]:

| | |
|----------------|----------|
| Firmicutes | 6.448889 |
| Proteobacteria | 7.401231 |
| Actinobacteria | 6.343880 |
| Bacteroidetes | 4.744932 |

dtype: float64

We can also filter according to the values in the **Series** :

```
In [13]: bacteria[bacteria>1000]
```

Out[13]:

| phylum | |
|----------------|------|
| Proteobacteria | 1638 |

Name: counts, dtype: int64

A **Series** can be thought of as an ordered key-value store. In fact, we can create one from a **dict** :

```
In [14]: bacteria_dict = { 'Firmicutes': 632,
                          'Proteobacteria': 1638,
                          'Actinobacteria': 569,
                          'Bacteroidetes': 115 }
pd.Series(bacteria_dict)
```

```
Out[14]: Firmicutes      632
Proteobacteria  1638
Actinobacteria   569
Bacteroidetes   115
dtype: int64
```

Notice that the `Series` is created in key-sorted order.

If we pass a custom index to `Series`, it will select the corresponding values from the dict, and treat indices without corresponding values as missing. Pandas uses the `NaN` (not a number) type for missing values.

```
In [15]: bacteria2 = pd.Series(bacteria_dict,
                              index=['Cyanobacteria', 'Firmicutes',
                                    'Proteobacteria', 'Actinobacteria'])
bacteria2
```

```
Out[15]: Cyanobacteria      NaN
Firmicutes      632.0
Proteobacteria  1638.0
Actinobacteria   569.0
dtype: float64
```

```
In [16]: bacteria2.isnull()
```

```
Out[16]: Cyanobacteria      True
Firmicutes      False
Proteobacteria  False
Actinobacteria  False
dtype: bool
```

Critically, the labels are used to **align data** when used in operations with other Series objects:

```
In [17]: bacteria + bacteria2
```

```
Out[17]: Actinobacteria      1138.0
Bacteroidetes      NaN
Cyanobacteria      NaN
Firmicutes      1264.0
Proteobacteria     3276.0
dtype: float64
```

Contrast this with NumPy arrays, where arrays of the same length will combine values **element-wise**; adding Series combined values with the same label in the resulting series. Notice also that the missing values were propagated by addition.

DataFrame

Inevitably, we want to be able to store, view and manipulate data that is *multivariate*, where for every index there are multiple fields or columns of data, often of varying data type.

A `DataFrame` is a tabular data structure, encapsulating multiple series like columns in a spreadsheet. Data are stored internally as a 2-dimensional object, but the `DataFrame` allows us to represent and manipulate higher-dimensional data.

```
In [13]: data = pd.DataFrame({'value':[632, 1638, 569, 115, 433, 1130, 754, 555],  
                             'patient':[1, 1, 1, 1, 2, 2, 2, 2],  
                             'phylum':['Firmicutes', 'Proteobacteria', 'Actinobacteria',  
                                         'Bacteroidetes', 'Firmicutes', 'Proteobacteria', 'Actinobacteria', 'Bacteroidetes']})  
data
```

```
Out[13]:
```

| | value | patient | phylum |
|---|-------|---------|----------------|
| 0 | 632 | 1 | Firmicutes |
| 1 | 1638 | 1 | Proteobacteria |
| 2 | 569 | 1 | Actinobacteria |
| 3 | 115 | 1 | Bacteroidetes |
| 4 | 433 | 2 | Firmicutes |
| 5 | 1130 | 2 | Proteobacteria |
| 6 | 754 | 2 | Actinobacteria |
| 7 | 555 | 2 | Bacteroidetes |

Notice the `DataFrame` is sorted by column name. We can change the order by indexing them in the order we desire:

```
In [14]: data[['phylum', 'value', 'patient']]
```

```
Out[14]:
```

| | phylum | value | patient |
|---|----------------|-------|---------|
| 0 | Firmicutes | 632 | 1 |
| 1 | Proteobacteria | 1638 | 1 |
| 2 | Actinobacteria | 569 | 1 |
| 3 | Bacteroidetes | 115 | 1 |
| 4 | Firmicutes | 433 | 2 |
| 5 | Proteobacteria | 1130 | 2 |
| 6 | Actinobacteria | 754 | 2 |
| 7 | Bacteroidetes | 555 | 2 |

A `DataFrame` has a second index, representing the columns:

```
In [41]: data.index
```

```
Out[41]: RangeIndex(start=0, stop=8, step=1)
```

The `dtypes` attribute reveals the data type for each column in our `DataFrame`.

- `int64` is numeric integer values
- `object` strings (letters and numbers)
- `float64` floating-point values

```
In [16]: data.dtypes
```

```
Out[16]: value      int64  
patient    int64
```

```
phylum      object
dtype: object
```

If we wish to access columns, we can do so either by dict-like indexing or by attribute:

```
In [17]: data['patient']
```

```
Out[17]:
```

| | |
|---|---|
| 0 | 1 |
| 1 | 1 |
| 2 | 1 |
| 3 | 1 |
| 4 | 2 |
| 5 | 2 |
| 6 | 2 |
| 7 | 2 |

Name: patient, dtype: int64

```
In [27]: x = data.value
         a = 1
```

```

/Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/_pydevd_bundle/pydevd_utils.py:606: FutureWarning: iteritems is deprecated and will be removed in a future version. Use .items instead.
    for item in s.iteritems():
/Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/_pydevd_bundle/pydevd_utils.py:606: FutureWarning: iteritems is deprecated and will be removed in a future version. Use .items instead.
    for item in s.iteritems():
/Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/_pydevd_bundle/pydevd_utils.py:606: FutureWarning: iteritems is deprecated and will be removed in a future version. Use .items instead.
    for item in s.iteritems():
/Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/_pydevd_bundle/pydevd_utils.py:606: FutureWarning: iteritems is deprecated and will be removed in a future version. Use .items instead.
    for item in s.iteritems():
/Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/_pydevd_bundle/pydevd_utils.py:606: FutureWarning: iteritems is deprecated and will be removed in a future version. Use .items instead.
    for item in s.iteritems():
-----
KeyboardInterrupt                                Traceback (most recent call last)
Cell In[27], line 2
      1 x = data.value
----> 2 a = 1

Cell In[27], line 2
      1 x = data.value
----> 2 a = 1

File _pydevd_bundle/pydevd_cython_darwin_39_64.pyx:1179, in _pydevd_bundle.pydevd_cython_darwin_39_64.SafeCallWrapper.__call__()

File _pydevd_bundle/pydevd_cython_darwin_39_64.pyx:620, in _pydevd_bundle.pydevd_cython_darwin_39_64.PyDBFrame.trace_dispatch()

File _pydevd_bundle/pydevd_cython_darwin_39_64.pyx:929, in _pydevd_bundle.pydevd_cython_darwin_39_64.PyDBFrame.trace_dispatch()

File _pydevd_bundle/pydevd_cython_darwin_39_64.pyx:920, in _pydevd_bundle.pydevd_cython_darwin_39_64.PyDBFrame.trace_dispatch()

```

```
File _pydevd_bundle/pydevd_cython_darwin_39_64.pyx:317, in _pydevd_bundle.pydevd_cython_darwin_39_64.PyDBFrame.do_wait_suspend()
```

```
File /Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/pydevd.py:1160, in PyDB.do_wait_suspend(self, thread, frame, event, arg, send_suspend_message, is_unhandled_exception)
```

```
1157         from_this_thread.append(frame_id)
1159 with self._threads_suspended_single_notification.notify_thread_suspended(thread_id, stop_reason):
-> 1160     self._do_wait_suspend(thread, frame, event, arg, suspend_type, from_this_thread)
```

```
File /Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/pydevd.py:1175, in PyDB._do_wait_suspend(self, thread, frame, event, arg, suspend_type, from_this_thread)
```

```
1172         self._call_mpl_hook()
1174         self.process_internal_commands()
-> 1175         time.sleep(0.01)
1177 self.cancel_async_evaluation(get_current_thread_id(thread), str(id(frame)))
1179 # process any stepping instructions
```

KeyboardInterrupt:

```
In [29]: x = data[['value']]
a = 1
```

```
In [30]: type(data[['value']])
```

```
Out[30]: pandas.core.frame.DataFrame
```

```
In [31]: type(data.value)
```

```
Out[31]: pandas.core.series.Series
```

Notice this is different than with `Series`, where dict-like indexing retrieved a particular element (row).

If we want access to a row in a `DataFrame`, we index its `loc` attribute.

```
In [27]: data.loc[3]
```

```
Out[27]: value          115
patient          1
phylum  Bacteroidetes
Name: 3, dtype: object
```

Exercise 1

Find the size of the first dimension of `data`.

```
In [28]: data.shape[0]
```

```
Out[28]: 8
```

An alternative way of initializing a `DataFrame` is with a list of dicts:

```
In [29]: data = pd.DataFrame([{'patient': 1, 'phylum': 'Firmicutes', 'value': 632},
                             {'patient': 1, 'phylum': 'Proteobacteria', 'value': 1638},
                             {'patient': 1, 'phylum': 'Actinobacteria', 'value': 569},
                             {'patient': 1, 'phylum': 'Bacteroidetes', 'value': 115},
                             {'patient': 2, 'phylum': 'Firmicutes', 'value': 433},
                             {'patient': 2, 'phylum': 'Proteobacteria', 'value': 1130},
```

```
{'patient': 2, 'phylum': 'Actinobacteria', 'value': 754},  
{'patient': 2, 'phylum': 'Bacteroidetes', 'value': 555}]]  
data
```

```
Out[29]:
```

| | patient | phylum | value |
|---|---------|----------------|-------|
| 0 | 1 | Firmicutes | 632 |
| 1 | 1 | Proteobacteria | 1638 |
| 2 | 1 | Actinobacteria | 569 |
| 3 | 1 | Bacteroidetes | 115 |
| 4 | 2 | Firmicutes | 433 |
| 5 | 2 | Proteobacteria | 1130 |
| 6 | 2 | Actinobacteria | 754 |
| 7 | 2 | Bacteroidetes | 555 |

It's important to note that the `Series` returned when a `DataFrame` is indexed is merely a **view** on the `DataFrame`, and not a copy of the data itself. So you must be cautious when manipulating this data:

```
In [30]: vals = data.value  
vals
```

```
Out[30]:
```

| | |
|---|------|
| 0 | 632 |
| 1 | 1638 |
| 2 | 569 |
| 3 | 115 |
| 4 | 433 |
| 5 | 1130 |
| 6 | 754 |
| 7 | 555 |

Name: value, dtype: int64

```
In [31]: vals[5] = 0  
data.value
```

```
Out[31]:
```

| | |
|---|------|
| 0 | 632 |
| 1 | 1638 |
| 2 | 569 |
| 3 | 115 |
| 4 | 433 |
| 5 | 0 |
| 6 | 754 |
| 7 | 555 |

Name: value, dtype: int64

If we plan on modifying an extracted Series, it's a good idea to make a copy.

```
In [32]: vals = data.value.copy()  
vals[5] = 1000  
data.value
```

```
Out[32]:
```

| | |
|---|------|
| 0 | 632 |
| 1 | 1638 |
| 2 | 569 |
| 3 | 115 |
| 4 | 433 |
| 5 | 0 |
| 6 | 754 |


```
7      555  
Name: value, dtype: int64
```

We can create or modify columns by assignment:

```
In [33]: data.value[[3,4,6]] = [14, 21, 5]  
data
```

```
Out[33]:
```

| | patient | phylum | value |
|---|---------|----------------|-------|
| 0 | 1 | Firmicutes | 632 |
| 1 | 1 | Proteobacteria | 1638 |
| 2 | 1 | Actinobacteria | 569 |
| 3 | 1 | Bacteroidetes | 14 |
| 4 | 2 | Firmicutes | 21 |
| 5 | 2 | Proteobacteria | 0 |
| 6 | 2 | Actinobacteria | 5 |
| 7 | 2 | Bacteroidetes | 555 |

```
In [34]: data['year'] = 2013  
data
```

```
Out[34]:
```

| | patient | phylum | value | year |
|---|---------|----------------|-------|------|
| 0 | 1 | Firmicutes | 632 | 2013 |
| 1 | 1 | Proteobacteria | 1638 | 2013 |
| 2 | 1 | Actinobacteria | 569 | 2013 |
| 3 | 1 | Bacteroidetes | 14 | 2013 |
| 4 | 2 | Firmicutes | 21 | 2013 |
| 5 | 2 | Proteobacteria | 0 | 2013 |
| 6 | 2 | Actinobacteria | 5 | 2013 |
| 7 | 2 | Bacteroidetes | 555 | 2013 |

But note we cannot use the attribute indexing method to add a new column:

```
In [35]: data.treatment = 1  
data
```

```
Out[35]:
```

| | patient | phylum | value | year |
|---|---------|----------------|-------|------|
| 0 | 1 | Firmicutes | 632 | 2013 |
| 1 | 1 | Proteobacteria | 1638 | 2013 |
| 2 | 1 | Actinobacteria | 569 | 2013 |
| 3 | 1 | Bacteroidetes | 14 | 2013 |
| 4 | 2 | Firmicutes | 21 | 2013 |
| 5 | 2 | Proteobacteria | 0 | 2013 |
| 6 | 2 | Actinobacteria | 5 | 2013 |
| 7 | 2 | Bacteroidetes | 555 | 2013 |

```
In [36]: data.treatment
```

```
Out[36]: 1
```

```
In [32]: data[[name.endswith('bacteria') for name in data.phylum] and data.value>1000]  
data[data['phylum'].apply(lambda x: x.endswith('bacteria')) & data['value'].apply(lambda
```

```
Out[32]:
```

| | value | patient | phylum |
|---|-------|---------|----------------|
| 1 | 1638 | 1 | Proteobacteria |
| 5 | 1130 | 2 | Proteobacteria |

Exercise 2

From the `data` table above, return all rows for which the phylum name ends in "bacteria" and the value is greater than 1000.

```
In [33]: data[data['phylum'].apply(lambda x: x.endswith('bacteria')) & data['value'].apply(lambda
```

```
Out[33]:
```

| | value | patient | phylum |
|---|-------|---------|----------------|
| 1 | 1638 | 1 | Proteobacteria |
| 5 | 1130 | 2 | Proteobacteria |

Specifying a `Series` as a new column causes its values to be added according to the `DataFrame`'s index:

```
In [35]: treatment = pd.Series([0]*4 + [1]*2)  
treatment
```

```
Out[35]:
```

| | |
|---|---|
| 0 | 0 |
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 4 | 1 |
| 5 | 1 |

dtype: int64

```
In [36]: data['treatment'] = treatment  
data
```

```
Out[36]:
```

| | value | patient | phylum | treatment |
|---|-------|---------|----------------|-----------|
| 0 | 632 | 1 | Firmicutes | 0.0 |
| 1 | 1638 | 1 | Proteobacteria | 0.0 |
| 2 | 569 | 1 | Actinobacteria | 0.0 |
| 3 | 115 | 1 | Bacteroidetes | 0.0 |
| 4 | 433 | 2 | Firmicutes | 1.0 |
| 5 | 1130 | 2 | Proteobacteria | 1.0 |
| 6 | 754 | 2 | Actinobacteria | NaN |
| 7 | 555 | 2 | Bacteroidetes | NaN |

We can use the `drop` method to remove rows or columns, which by **default drops rows**. We can be

explicit by using the `axis` argument:

```
In [37]: data_nomonth = data.drop('treatment', axis=1)
data_nomonth
```

```
Out[37]:
```

| | value | patient | phylum |
|---|-------|---------|----------------|
| 0 | 632 | 1 | Firmicutes |
| 1 | 1638 | 1 | Proteobacteria |
| 2 | 569 | 1 | Actinobacteria |
| 3 | 115 | 1 | Bacteroidetes |
| 4 | 433 | 2 | Firmicutes |
| 5 | 1130 | 2 | Proteobacteria |
| 6 | 754 | 2 | Actinobacteria |
| 7 | 555 | 2 | Bacteroidetes |

```
In [38]: data_nomonth.drop(0)
```

```
Out[38]:
```

| | value | patient | phylum |
|---|-------|---------|----------------|
| 1 | 1638 | 1 | Proteobacteria |
| 2 | 569 | 1 | Actinobacteria |
| 3 | 115 | 1 | Bacteroidetes |
| 4 | 433 | 2 | Firmicutes |
| 5 | 1130 | 2 | Proteobacteria |
| 6 | 754 | 2 | Actinobacteria |
| 7 | 555 | 2 | Bacteroidetes |

The dtype is automatically chosen to be as general as needed to accommodate all the columns.

```
In [39]: df = pd.DataFrame({'foo': [1,2,3], 'bar':[0.4, -1.0, 4.5]})
df.values
```

```
Out[39]: array([[ 1. ,  0.4],
               [ 2. , -1. ],
               [ 3. ,  4.5]])
```

```
In [40]: df.values.dtype
```

```
Out[40]: dtype('float64')
```

Index objects are **immutable**:

```
In [44]: #data.index[0] = 15 #that doesn't work
```

This is so that **Index** objects can be shared between data structures without fear that they will be changed.

```
In [45]: bacteria2.index = bacteria.index
```

```
In [46]: bacteria2
```

```
Out[46]: phylum
          Firmicutes      NaN
          Proteobacteria  632.0
          Actinobacteria  1638.0
          Bacteroidetes   569.0
dtype: float64
```

Importing data

A key, but often under-appreciated, step in data analysis is importing the data that we wish to analyze. Though it is easy to load basic data structures into Python using built-in tools or those provided by packages like NumPy, it is non-trivial to import structured data well, and to easily convert this input into a robust data structure:

```
genes = np.loadtxt("genes.csv", delimiter=";", dtype=[('gene', '<S10'),
               ('value', '<f4')])
```

Pandas provides a convenient set of functions for importing tabular data from a number of formats directly into a `DataFrame` object. These functions include a slew of options to perform type inference, indexing, parsing, iterating and cleaning automatically as data are imported.

Let's start with some more bacteria data, stored in csv format.

```
In [43]: !head Data/microbiome.csv

Taxon,Patient,Group,Tissue,Stool
Firmicutes,1,0,136,4182
Firmicutes,2,1,1174,703
Firmicutes,3,0,408,3946
Firmicutes,4,1,831,8605
Firmicutes,5,0,693,50
Firmicutes,6,1,718,717
Firmicutes,7,0,173,33
Firmicutes,8,1,228,80
Firmicutes,9,0,162,3196
```

This table can be read into a DataFrame using `read_csv`:

```
In [46]: mb = pd.read_csv("Data/microbiome.csv", sep=',')
         mb.head()
```

```
Out[46]:
```

| | Taxon | Patient | Group | Tissue | Stool |
|---|------------|---------|-------|--------|-------|
| 0 | Firmicutes | 1 | 0 | 136 | 4182 |
| 1 | Firmicutes | 2 | 1 | 1174 | 703 |
| 2 | Firmicutes | 3 | 0 | 408 | 3946 |
| 3 | Firmicutes | 4 | 1 | 831 | 8605 |
| 4 | Firmicutes | 5 | 0 | 693 | 50 |

Notice that `read_csv` automatically considered the first row in the file to be a header row.

We can override default behavior by customizing some the arguments, like `header`, `names` or `index_col`.

```
In [49]: pd.read_csv("Data/microbiome.csv", header=None).head()
```

```
Out[49]:
```

| | 0 | 1 | 2 | 3 | 4 |
|---|------------|---------|-------|--------|-------|
| 0 | Taxon | Patient | Group | Tissue | Stool |
| 1 | Firmicutes | 1 | 0 | 136 | 4182 |
| 2 | Firmicutes | 2 | 1 | 1174 | 703 |
| 3 | Firmicutes | 3 | 0 | 408 | 3946 |
| 4 | Firmicutes | 4 | 1 | 831 | 8605 |

`read_csv` is just a convenience function for `read_table`, since csv is such a common format:

```
In [50]: mb = pd.read_csv("Data/microbiome.csv", sep=',')
```

The `sep` argument can be customized as needed to accomodate arbitrary separators. For example, we can use a regular expression to define a variable amount of whitespace, which is unfortunately very common in some data formats:

```
sep='\s+'
```

For a more useful index, we can specify the first two columns, which together provide a unique index to the data.

```
In [51]: mb = pd.read_csv("Data/microbiome.csv", index_col=['Patient', 'Taxon'])
mb.head()
```

```
Out[51]:
```

| | | Group | Tissue | Stool |
|---|------------|-------|--------|-------|
| | Patient | Taxon | | |
| 1 | Firmicutes | 0 | 136 | 4182 |
| 2 | Firmicutes | 1 | 1174 | 703 |
| 3 | Firmicutes | 0 | 408 | 3946 |
| 4 | Firmicutes | 1 | 831 | 8605 |
| 5 | Firmicutes | 0 | 693 | 50 |

This is called a *hierarchical* index, which we will revisit later in the section.

If we have sections of data that we do not wish to import (for example, known bad data), we can populate the `skiprows` argument:

```
In [52]: pd.read_csv("Data/microbiome.csv", skiprows=[3,4,6]).head()
```

```
Out[52]:
```

| | Taxon | Patient | Group | Tissue | Stool |
|---|------------|---------|-------|--------|-------|
| 0 | Firmicutes | 1 | 0 | 136 | 4182 |
| 1 | Firmicutes | 2 | 1 | 1174 | 703 |
| 2 | Firmicutes | 5 | 0 | 693 | 50 |
| 3 | Firmicutes | 7 | 0 | 173 | 33 |
| 4 | Firmicutes | 8 | 1 | 228 | 80 |

If we only want to import a small number of rows from, say, a very large data file we can use `nrows` :

```
In [53]: pd.read_csv("Data/microbiome.csv", nrows=4)
```

```
Out[53]:
```

| | Taxon | Patient | Group | Tissue | Stool |
|---|------------|---------|-------|--------|-------|
| 0 | Firmicutes | 1 | 0 | 136 | 4182 |
| 1 | Firmicutes | 2 | 1 | 1174 | 703 |
| 2 | Firmicutes | 3 | 0 | 408 | 3946 |
| 3 | Firmicutes | 4 | 1 | 831 | 8605 |

Missing Values

Most real-world data is incomplete, with values missing due to incomplete observation, data entry or transcription error, or other reasons. Pandas will automatically recognize and parse common missing data indicators, including `NA` and `NULL` .

```
In [45]: !head -n 10 Data/microbiome_missing.csv
```

```
Taxon,Patient,Tissue,Stool
Firmicutes,1,632,305
Firmicutes,2,136,4182
Firmicutes,3,,703
Firmicutes,4,408,3946
Firmicutes,5,831,8605
Firmicutes,6,693,50
Firmicutes,7,718,717
Firmicutes,8,173,33
Firmicutes,9,228,NA
```

```
In [55]: pd.read_csv("Data/microbiome_missing.csv").head(10)
```

```
Out[55]:
```

| | Taxon | Patient | Tissue | Stool |
|---|------------|---------|--------|--------|
| 0 | Firmicutes | 1 | 632 | 305.0 |
| 1 | Firmicutes | 2 | 136 | 4182.0 |
| 2 | Firmicutes | 3 | NaN | 703.0 |
| 3 | Firmicutes | 4 | 408 | 3946.0 |
| 4 | Firmicutes | 5 | 831 | 8605.0 |
| 5 | Firmicutes | 6 | 693 | 50.0 |
| 6 | Firmicutes | 7 | 718 | 717.0 |
| 7 | Firmicutes | 8 | 173 | 33.0 |
| 8 | Firmicutes | 9 | 228 | NaN |
| 9 | Firmicutes | 10 | 162 | 3196.0 |

Above, Pandas recognized `NA` and an empty field as missing data.

```
In [56]: pd.isnull(pd.read_csv("Data/microbiome_missing.csv")).head(10)
```

```
Out[56]:
```

| | Taxon | Patient | Tissue | Stool |
|---|-------|---------|--------|-------|
| 0 | False | False | False | False |

| | | | | |
|---|-------|-------|-------|-------|
| 1 | False | False | False | False |
| 2 | False | False | True | False |
| 3 | False | False | False | False |
| 4 | False | False | False | False |
| 5 | False | False | False | False |
| 6 | False | False | False | False |
| 7 | False | False | False | False |
| 8 | False | False | False | True |
| 9 | False | False | False | False |

Unfortunately, there will sometimes be inconsistencies in the conventions for missing data. In this example, there is a question mark "?" and a large negative number where there should have been a positive integer. We can specify additional symbols with the `na_values` argument:

```
In [57]: pd.read_csv("Data/microbiome_missing.csv", na_values=['?', -99999]).head(10)
```

```
Out[57]:
```

| | Taxon | Patient | Tissue | Stool |
|--|-------|---------|--------|-------|
|--|-------|---------|--------|-------|

| | | | | |
|---|------------|----|-------|--------|
| 0 | Firmicutes | 1 | 632.0 | 305.0 |
| 1 | Firmicutes | 2 | 136.0 | 4182.0 |
| 2 | Firmicutes | 3 | NaN | 703.0 |
| 3 | Firmicutes | 4 | 408.0 | 3946.0 |
| 4 | Firmicutes | 5 | 831.0 | 8605.0 |
| 5 | Firmicutes | 6 | 693.0 | 50.0 |
| 6 | Firmicutes | 7 | 718.0 | 717.0 |
| 7 | Firmicutes | 8 | 173.0 | 33.0 |
| 8 | Firmicutes | 9 | 228.0 | NaN |
| 9 | Firmicutes | 10 | 162.0 | 3196.0 |

These can be specified on a column-wise basis using an appropriate dict as the argument for `na_values`.

Microsoft Excel

Since so much financial and scientific data ends up in Excel spreadsheets (regrettably), Pandas' ability to directly import Excel spreadsheets is valuable. This support is contingent on having one or two dependencies (depending on what version of Excel file is being imported) installed: `xlrd` and `openpyxl` (these may be installed with either `pip` or `easy_install`).

The `read_excel` convenience function in pandas imports a specific sheet from an Excel file

```
In [58]: mb = pd.read_excel('Data/microbiome_MID2.xls', sheet_name='Sheet 1', header=None)
mb.head()
```

```
Out[58]:
```

| | 0 | 1 |
|---|---|---|
| 0 | Archaea "Crenarchaeota" Thermoprotei Acidiloba... | 2 |

| | | |
|---|---|----|
| 1 | Archaea "Crenarchaeota" Thermoprotei Acidiloba... | 14 |
| 2 | Archaea "Crenarchaeota" Thermoprotei Desulfuro... | 23 |
| 3 | Archaea "Crenarchaeota" Thermoprotei Desulfuro... | 1 |
| 4 | Archaea "Crenarchaeota" Thermoprotei Desulfuro... | 2 |

There are several other data formats that can be imported into Python and converted into `DataFrames`, with the help of built-in or third-party libraries. These include JSON, XML, HDF5, relational and non-relational databases, and various web APIs. These are beyond the scope of this tutorial, but are covered in [Python for Data Analysis](#).

Indexing and Selection

This section introduces the new user to the key functionality of Pandas that is required to use the software effectively.

For some variety, we will leave our digestive tract bacteria behind and employ some baseball data.

```
In [59]: baseball = pd.read_csv("Data/baseball.csv", index_col='id')
baseball.head()
```

```
Out[59]:
```

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb | so | ibb | hbp | sh |
|--------------|-----------|------|-------|------|----|----|----|---|----|-----|-----|-----|-----|-----|----|-----|-----|-----|-----|
| id | | | | | | | | | | | | | | | | | | | |
| 88641 | womacto01 | 2006 | 2 | CHN | NL | 19 | 50 | 6 | 14 | 1 | ... | 2.0 | 1.0 | 1.0 | 4 | 4.0 | 0.0 | 0.0 | 3.0 |
| 88643 | schilcu01 | 2006 | 1 | BOS | AL | 31 | 2 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 1.0 | 0.0 | 0.0 | 0.0 |
| 88645 | myersmi01 | 2006 | 1 | NYA | AL | 62 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 88649 | helliri01 | 2006 | 1 | MIL | NL | 20 | 3 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 2.0 | 0.0 | 0.0 | 0.0 |
| 88650 | johnsra05 | 2006 | 1 | NYA | AL | 33 | 6 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 4.0 | 0.0 | 0.0 | 0.0 |

5 rows × 22 columns

Notice that we specified the `id` column as the index, since it appears to be a unique identifier. We could try to create a unique index ourselves by combining `player` and `year`:

```
In [60]: baseball.index.is_unique
```

```
Out[60]: True
```

```
In [61]: player_id = baseball.player + baseball.year.astype(str)
baseball_newind = baseball.copy()
baseball_newind.index = player_id
baseball_newind.head()
```

```
Out[61]:
```

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb | so | ibb |
|----------------------|-----------|------|-------|------|----|----|----|---|----|-----|-----|-----|-----|-----|----|-----|-----|
| womacto012006 | womacto01 | 2006 | 2 | CHN | NL | 19 | 50 | 6 | 14 | 1 | ... | 2.0 | 1.0 | 1.0 | 4 | 4.0 | 0.0 |
| schilcu012006 | schilcu01 | 2006 | 1 | BOS | AL | 31 | 2 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 1.0 | 0.0 |
| myersmi012006 | myersmi01 | 2006 | 1 | NYA | AL | 62 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 0.0 | 0.0 |
| helliri012006 | helliri01 | 2006 | 1 | MIL | NL | 20 | 3 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 2.0 | 0.0 |
| johnsra052006 | johnsra05 | 2006 | 1 | NYA | AL | 33 | 6 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 4.0 | 0.0 |

5 rows × 22 columns

This looks okay, but let's check:

```
In [62]: baseball_newind.index.is_unique
```

```
Out[62]: False
```

So, indices need not be unique. Our choice is not unique because some players change teams within years.

The most important consequence of a non-unique index is that indexing by label will return multiple values for some labels:

```
In [63]: baseball_newind.loc['wickmbo012007']
```

```
Out[63]:
```

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb | so | ibb | ht |
|----------------------|-----------|------|-------|------|----|----|----|---|---|-----|-----|-----|-----|-----|----|-----|-----|----|
| wickmbo012007 | wickmbo01 | 2007 | 2 | ARI | NL | 8 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 0.0 | 0.0 | 0 |
| wickmbo012007 | wickmbo01 | 2007 | 1 | ATL | NL | 47 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 0.0 | 0.0 | 0 |

2 rows × 22 columns

We will learn more about indexing below.

We can create a truly unique index by combining `player`, `team` and `year`:

```
In [64]: player_unique = baseball.player + baseball.team + baseball.year.astype(str)
baseball_newind = baseball.copy()
baseball_newind.index = player_unique
baseball_newind.head()
```

```
Out[64]:
```

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb | so | i |
|-------------------------|-----------|------|-------|------|----|----|----|---|----|-----|-----|-----|-----|-----|----|-----|---|
| womacto01CHN2006 | womacto01 | 2006 | 2 | CHN | NL | 19 | 50 | 6 | 14 | 1 | ... | 2.0 | 1.0 | 1.0 | 4 | 4.0 | 0 |
| schilcu01BOS2006 | schilcu01 | 2006 | 1 | BOS | AL | 31 | 2 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 1.0 | 0 |
| myersmi01NYA2006 | myersmi01 | 2006 | 1 | NYA | AL | 62 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 0.0 | 0 |
| helliri01MIL2006 | helliri01 | 2006 | 1 | MIL | NL | 20 | 3 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 2.0 | 0 |
| johnsra05NYA2006 | johnsra05 | 2006 | 1 | NYA | AL | 33 | 6 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 4.0 | 0 |

5 rows × 22 columns

```
In [65]: baseball_newind.index.is_unique
```

```
Out[65]: True
```

We can create meaningful indices more easily using a hierarchical index; for now, we will stick with the numeric `id` field as our index.

Exercise 3

Build another unique index without using the same columns as `player_unique`.

```
In [66]: player_unique = baseball.team + baseball.year.astype(str) + baseball.player
player_unique.is_unique
```

Out[66]: True

Indexing works analogously to indexing in NumPy arrays, except we can use the labels in the **Index** object to extract values in addition to arrays of integers.

```
In [67]: # Numpy-style indexing
baseball_newind[:3]
```

Out[67]:

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb | so | it |
|------------------|-----------|------|-------|------|----|----|----|---|----|-----|-----|-----|-----|-----|----|-----|----|
| womacto01CHN2006 | womacto01 | 2006 | 2 | CHN | NL | 19 | 50 | 6 | 14 | 1 | ... | 2.0 | 1.0 | 1.0 | 4 | 4.0 | 0 |
| schilcu01BOS2006 | schilcu01 | 2006 | 1 | BOS | AL | 31 | 2 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 1.0 | 0 |
| myersmi01NYA2006 | myersmi01 | 2006 | 1 | NYA | AL | 62 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 0.0 | 0 |

3 rows × 22 columns

We can also slice with data labels, since they have an intrinsic order within the Index:

```
In [68]: baseball_newind['womacto01CHN2006':'gonzalu01ARI2006']
```

Out[68]:

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb | |
|------------------|-----------|------|-------|------|----|-----|-----|----|-----|-----|-----|------|-----|-----|----|--|
| womacto01CHN2006 | womacto01 | 2006 | 2 | CHN | NL | 19 | 50 | 6 | 14 | 1 | ... | 2.0 | 1.0 | 1.0 | 4 | |
| schilcu01BOS2006 | schilcu01 | 2006 | 1 | BOS | AL | 31 | 2 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | |
| myersmi01NYA2006 | myersmi01 | 2006 | 1 | NYA | AL | 62 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | |
| helliri01MIL2006 | helliri01 | 2006 | 1 | MIL | NL | 20 | 3 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | |
| johnsra05NYA2006 | johnsra05 | 2006 | 1 | NYA | AL | 33 | 6 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | |
| finlest01SFN2006 | finlest01 | 2006 | 1 | SFN | NL | 139 | 426 | 66 | 105 | 21 | ... | 40.0 | 7.0 | 0.0 | 46 | |
| gonzalu01ARI2006 | gonzalu01 | 2006 | 1 | ARI | NL | 153 | 586 | 93 | 159 | 52 | ... | 73.0 | 0.0 | 1.0 | 69 | |

7 rows × 22 columns

```
In [69]: baseball_newind['womacto01CHN2006':'gonzalu01ARI2006'] = 5
baseball_newind.head(10)
```

Out[69]:

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb | so | it |
|------------------|-----------|------|-------|------|----|----|----|---|----|-----|-----|-----|-----|-----|----|------|----|
| womacto01CHN2006 | | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | ... | 5.0 | 5.0 | 5.0 | 5 | 5.0 | 5 |
| schilcu01BOS2006 | | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | ... | 5.0 | 5.0 | 5.0 | 5 | 5.0 | 5 |
| myersmi01NYA2006 | | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | ... | 5.0 | 5.0 | 5.0 | 5 | 5.0 | 5 |
| helliri01MIL2006 | | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | ... | 5.0 | 5.0 | 5.0 | 5 | 5.0 | 5 |
| johnsra05NYA2006 | | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | ... | 5.0 | 5.0 | 5.0 | 5 | 5.0 | 5 |
| finlest01SFN2006 | | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | ... | 5.0 | 5.0 | 5.0 | 5 | 5.0 | 5 |
| gonzalu01ARI2006 | | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | ... | 5.0 | 5.0 | 5.0 | 5 | 5.0 | 5 |
| seleaa01LAN2006 | seleaa01 | 2006 | 1 | LAN | NL | 28 | 26 | 2 | 5 | 1 | ... | 0.0 | 0.0 | 0.0 | 1 | 7.0 | 0 |
| francju01ATL2007 | francju01 | 2007 | 2 | ATL | NL | 15 | 40 | 1 | 10 | 3 | ... | 8.0 | 0.0 | 0.0 | 4 | 10.0 | 1 |

10 rows × 22 columns

For a more concise (and readable) syntax, we can use the new `query` method to perform selection on a `DataFrame`. Instead of having to type the fully-specified column, we can simply pass a string that describes what to select. The query above is then simply:

```
In [70]: baseball_newind.query('ab > 500')
```

```
Out[70]:
```

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb |
|-------------------------|-----------|------|-------|------|----|-----|-----|----|-----|-----|-----|------|------|-----|----|
| vizquom01SFN2007 | vizquom01 | 2007 | 1 | SFN | NL | 145 | 513 | 54 | 126 | 18 | ... | 51.0 | 14.0 | 6.0 | 44 |
| thomafr04TOR2007 | thomafr04 | 2007 | 1 | TOR | AL | 155 | 531 | 63 | 147 | 30 | ... | 95.0 | 0.0 | 0.0 | 81 |
| rodriiv01DET2007 | rodriiv01 | 2007 | 1 | DET | AL | 129 | 502 | 50 | 141 | 31 | ... | 63.0 | 2.0 | 2.0 | 9 |
| griffke02CIN2007 | griffke02 | 2007 | 1 | CIN | NL | 144 | 528 | 78 | 146 | 24 | ... | 93.0 | 6.0 | 1.0 | 85 |
| delgaca01NYN2007 | delgaca01 | 2007 | 1 | NYN | NL | 139 | 538 | 71 | 139 | 30 | ... | 87.0 | 4.0 | 0.0 | 52 |
| biggicr01HOU2007 | biggicr01 | 2007 | 1 | HOU | NL | 141 | 517 | 68 | 130 | 31 | ... | 50.0 | 4.0 | 3.0 | 23 |

6 rows × 22 columns

The `DataFrame.index` and `DataFrame.columns` are placed in the query namespace by default. If you want to refer to a variable in the current namespace, you can prefix the variable with `@`:

```
In [71]: min_ab = 500
```

```
In [72]: baseball_newind.query('ab > @min_ab')
```

```
Out[72]:
```

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb |
|-------------------------|-----------|------|-------|------|----|-----|-----|----|-----|-----|-----|------|------|-----|----|
| vizquom01SFN2007 | vizquom01 | 2007 | 1 | SFN | NL | 145 | 513 | 54 | 126 | 18 | ... | 51.0 | 14.0 | 6.0 | 44 |
| thomafr04TOR2007 | thomafr04 | 2007 | 1 | TOR | AL | 155 | 531 | 63 | 147 | 30 | ... | 95.0 | 0.0 | 0.0 | 81 |
| rodriiv01DET2007 | rodriiv01 | 2007 | 1 | DET | AL | 129 | 502 | 50 | 141 | 31 | ... | 63.0 | 2.0 | 2.0 | 9 |
| griffke02CIN2007 | griffke02 | 2007 | 1 | CIN | NL | 144 | 528 | 78 | 146 | 24 | ... | 93.0 | 6.0 | 1.0 | 85 |
| delgaca01NYN2007 | delgaca01 | 2007 | 1 | NYN | NL | 139 | 538 | 71 | 139 | 30 | ... | 87.0 | 4.0 | 0.0 | 52 |
| biggicr01HOU2007 | biggicr01 | 2007 | 1 | HOU | NL | 141 | 517 | 68 | 130 | 31 | ... | 50.0 | 4.0 | 3.0 | 23 |

6 rows × 22 columns

The indexing field `loc` allows us to select subsets of rows and columns in an intuitive way:

```
In [73]: baseball_newind.loc['gonzalu01ARI2006', ['h', 'X2b', 'X3b', 'hr']]
```

```
Out[73]:
```

| | |
|-----|---|
| h | 5 |
| X2b | 5 |
| X3b | 5 |
| hr | 5 |

Name: gonzalu01ARI2006, dtype: object

```
In [74]: baseball_newind.loc[:'myersmi01NYA2006', 'hr']
```

```
Out[74]:
```

| | |
|------------------|---|
| womacto01CHN2006 | 5 |
| schilcu01BOS2006 | 5 |

```
myersmi01NYA2006    5
Name: hr, dtype: int64
```

In addition to using `loc` to select rows and columns by **label**, pandas also allows indexing by **position** using the `iloc` attribute.

So, we can query rows and columns by absolute position, rather than by name:

```
In [75]: baseball_newind.iloc[:5, 5:8]
```

```
Out[75]:
```

| | g | ab | r |
|------------------|---|----|---|
| womacto01CHN2006 | 5 | 5 | 5 |
| schilcu01BOS2006 | 5 | 5 | 5 |
| myersmi01NYA2006 | 5 | 5 | 5 |
| helliri01MIL2006 | 5 | 5 | 5 |
| johnsra05NYA2006 | 5 | 5 | 5 |

Exercise 4

You can use the `isin` method query a DataFrame based upon a list of values as follows:

```
data['phylum'].isin(['Firmicutes', 'Bacteroidetes'])
```

Use `isin` on `baseball` to find all players that played for the Los Angeles Dodgers (LAN) or the San Francisco Giants (SFN). How many records contain these values?

```
In [76]: len(baseball[baseball.team.isin(['LAN', 'SFN'])])
```

```
Out[76]: 15
```

Hierarchical indexing

In the baseball example, it was necessary to combine 3 fields to obtain a unique index that was not simply an integer value. A more elegant way to have done this would be to create a hierarchical index from the three fields.

```
In [77]: baseball_h = baseball.set_index(['year', 'team', 'player'])
baseball_h.head(10)
```

```
Out[77]:
```

| | | | stint | lg | g | ab | r | h | X2b | X3b | hr | rbi | sb | cs | bb | so | ibb | hbp |
|------|------|-----------|--------|----|-----|-----|----|-----|-----|-----|----|------|-----|-----|----|------|-----|-----|
| | year | team | player | | | | | | | | | | | | | | | |
| 2006 | CHN | womacto01 | 2 | NL | 19 | 50 | 6 | 14 | 1 | 0 | 1 | 2.0 | 1.0 | 1.0 | 4 | 4.0 | 0.0 | 0.0 |
| | BOS | schilcu01 | 1 | AL | 31 | 2 | 0 | 1 | 0 | 0 | 0 | 0.0 | 0.0 | 0.0 | 0 | 1.0 | 0.0 | 0.0 |
| | NYA | myersmi01 | 1 | AL | 62 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0.0 | 0.0 | 0 | 0.0 | 0.0 | 0.0 |
| | MIL | helliri01 | 1 | NL | 20 | 3 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0.0 | 0.0 | 0 | 2.0 | 0.0 | 0.0 |
| | NYA | johnsra05 | 1 | AL | 33 | 6 | 0 | 1 | 0 | 0 | 0 | 0.0 | 0.0 | 0.0 | 0 | 4.0 | 0.0 | 0.0 |
| | SFN | finlest01 | 1 | NL | 139 | 426 | 66 | 105 | 21 | 12 | 6 | 40.0 | 7.0 | 0.0 | 46 | 55.0 | 2.0 | 2.0 |

| | | | | | | | | | | | | | | | | | | |
|------|-----|-----------|---|----|-----|-----|----|-----|----|---|----|------|-----|-----|----|------|------|-----|
| | ARI | gonzalu01 | 1 | NL | 153 | 586 | 93 | 159 | 52 | 2 | 15 | 73.0 | 0.0 | 1.0 | 69 | 58.0 | 10.0 | 7.0 |
| | LAN | seleaa01 | 1 | NL | 28 | 26 | 2 | 5 | 1 | 0 | 0 | 0.0 | 0.0 | 0.0 | 1 | 7.0 | 0.0 | 0.0 |
| 2007 | ATL | francju01 | 2 | NL | 15 | 40 | 1 | 10 | 3 | 0 | 0 | 8.0 | 0.0 | 0.0 | 4 | 10.0 | 1.0 | 0.0 |
| | NYN | francju01 | 1 | NL | 40 | 50 | 7 | 10 | 0 | 0 | 1 | 8.0 | 2.0 | 1.0 | 10 | 13.0 | 0.0 | 0.0 |

```
In [78]: baseball_h.index.is_unique
```

Out[78]: True

Try using this hierarchical index to retrieve Julio Franco (`francju01`), who played for the Atlanta Braves (`ATL`) in 2007:

```
In [79]: baseball_h.loc[(2007, 'ATL', 'francju01')]
```

```
Out[79]: stint      2
lg          NL
g           15
ab          40
r            1
h           10
X2b         3
X3b         0
hr           0
rbi         8.0
sb           0.0
cs           0.0
bb           4
so          10.0
ibb          1.0
hbp          0.0
sh           0.0
sf           1.0
gidp         1.0
Name: (2007, ATL, francju01), dtype: object
```

Recall earlier we imported some microbiome data using two index columns. This created a 2-level hierarchical index:

```
In [80]: mb = pd.read_csv("Data/microbiome.csv", index_col=['Taxon','Patient'])
```

```
In [81]: mb.head(10)
```

```
Out[81]:
```

| | | Group | Tissue | Stool |
|--|------------|---------|--------|----------|
| | Taxon | Patient | | |
| | Firmicutes | 1 | 0 | 136 4182 |
| | | 2 | 1 | 1174 703 |
| | | 3 | 0 | 408 3946 |
| | | 4 | 1 | 831 8605 |
| | | 5 | 0 | 693 50 |
| | | 6 | 1 | 718 717 |
| | | 7 | 0 | 173 33 |
| | | 8 | 1 | 228 80 |
| | | 9 | 0 | 162 3196 |

With a hierarchical index, we can select subsets of the data based on a *partial* index:

```
In [82]: mb.loc['Proteobacteria']
```

```
Out[82]:
```

| | Group | Tissue | Stool |
|---------|-------|--------|-------|
| Patient | | | |
| 1 | 0 | 2469 | 1821 |
| 2 | 1 | 839 | 661 |
| 3 | 0 | 4414 | 18 |
| 4 | 1 | 12044 | 83 |
| 5 | 0 | 2310 | 12 |
| 6 | 1 | 3053 | 547 |
| 7 | 0 | 395 | 2174 |
| 8 | 1 | 2651 | 767 |
| 9 | 0 | 1195 | 76 |
| 10 | 1 | 6857 | 795 |
| 11 | 0 | 483 | 666 |
| 12 | 1 | 2950 | 3994 |
| 13 | 0 | 1541 | 816 |
| 14 | 1 | 1307 | 53 |

Hierarchical indices can be created on either or **both axes**. Here is a trivial example:

```
In [83]: frame = pd.DataFrame(np.arange(12).reshape(( 4, 3)),
                             index = [['a', 'a', 'b', 'b'], [1, 2, 1, 2]],
                             columns = [['Ohio', 'Ohio', 'Colorado'], ['Green', 'Red', 'Green']])

frame
```

```
Out[83]:
```

| | | Ohio | Colorado | |
|---|---|-------|----------|-------|
| | | Green | Red | Green |
| a | 1 | 0 | 1 | 2 |
| | 2 | 3 | 4 | 5 |
| b | 1 | 6 | 7 | 8 |
| | 2 | 9 | 10 | 11 |

If you want to get fancy, both the row and column indices themselves can be given names:

```
In [84]: frame.index.names = ['key1', 'key2']
         frame.columns.names = ['state', 'color']
         frame
```

```
Out[84]:
```

| | | | |
|-------|-------|----------|-------|
| state | Ohio | Colorado | |
| color | Green | Red | Green |

| key1 | key2 | | | |
|------|------|---|----|----|
| a | 1 | 0 | 1 | 2 |
| | 2 | 3 | 4 | 5 |
| b | 1 | 6 | 7 | 8 |
| | 2 | 9 | 10 | 11 |

Additionally, the order of the set of indices in a hierarchical `MultiIndex` can be changed by swapping them pairwise:

```
In [85]: frame.columns = frame.columns.swaplevel('state', 'color')
         frame
```

```
Out[85]:
```

| | color | Green | Red | Green |
|------|-------|-------|------|----------|
| | state | Ohio | Ohio | Colorado |
| key1 | key2 | | | |
| a | 1 | 0 | 1 | 2 |
| | 2 | 3 | 4 | 5 |
| b | 1 | 6 | 7 | 8 |
| | 2 | 9 | 10 | 11 |

Operations

`DataFrame` and `Series` objects allow for several operations to take place either on a single object, or between two or more objects.

For example, we can perform arithmetic on the elements of two objects, such as combining baseball statistics across years. First, let's (artificially) construct two `Series`, consisting of home runs hit in years 2006 and 2007, respectively:

```
In [86]: hr2006 = baseball.loc[baseball.year==2006, 'hr']
         hr2006.index = baseball.player[baseball.year==2006]

         hr2007 = baseball.loc[baseball.year==2007, 'hr']
         hr2007.index = baseball.player[baseball.year==2007]
```

```
In [87]: hr2007.head(10)
```

```
Out[87]:
```

| player | |
|-----------|----|
| francju01 | 0 |
| francju01 | 1 |
| zaungr01 | 10 |
| witasja01 | 0 |
| williwo02 | 1 |
| wickmbo01 | 0 |
| wickmbo01 | 0 |
| whitero02 | 4 |
| whiteri01 | 0 |
| welllda01 | 0 |

Name: hr, dtype: int64

Now, let's add them together, in hopes of getting 2-year home run totals:

```
In [88]: hr_total = hr2006 + hr2007
hr_total.head(20)
```

```
Out[88]: player
alomasa02      NaN
aloumo01       NaN
ausmubr01      NaN
benitar01      NaN
benitar01      NaN
biggicr01      NaN
bondsba01      NaN
cirilje01      NaN
cirilje01      NaN
claytro01      NaN
claytro01      NaN
clemero02      NaN
coninje01      NaN
coninje01      NaN
cormirh01      NaN
delgaca01      NaN
easleda01      NaN
edmonji01      NaN
embreal01      NaN
finlest01      7.0
Name: hr, dtype: float64
```

Pandas' data alignment places **NaN** values for labels that do not overlap in the two Series. In fact, there are only 6 players that occur in both years.

```
In [89]: hr_total[hr_total.notnull()]
```

```
Out[89]: player
finlest01      7.0
gonzalu01     30.0
johnsra05      0.0
myersmi01      0.0
schilcu01      0.0
seleaa01       0.0
Name: hr, dtype: float64
```

While we do want the operation to honor the data labels in this way, we probably do not want the missing values to be filled with **NaN**. We can use the **add** method to calculate player home run totals by using the **fill_value** argument to insert a zero for home runs where labels do not overlap:

```
In [90]: hr2007.add(hr2006, fill_value=0).head(20)
```

```
Out[90]: player
alomasa02      0.0
aloumo01     13.0
ausmubr01      3.0
benitar01      0.0
benitar01      0.0
biggicr01     10.0
bondsba01     28.0
cirilje01      0.0
cirilje01      2.0
claytro01      0.0
claytro01      1.0
clemero02      0.0
coninje01      0.0
coninje01      6.0
cormirh01      0.0
delgaca01     24.0
easleda01     10.0
```



```
edmonji01      12.0
embreal01       0.0
finlest01       7.0
Name: hr, dtype: float64
```

Operations can also be **broadcast** between rows or columns.

For example, if we subtract the maximum number of home runs hit from the `hr` column, we get how many fewer than the maximum were hit by each player:

```
In [91]: (baseball.hr - baseball.hr.max()).head(20)

Out[91]: id
88641    -34
88643    -35
88645    -35
88649    -35
88650    -35
88652    -29
88653    -20
88662    -35
89177    -35
89178    -34
89330    -25
89333    -35
89334    -34
89335    -35
89336    -35
89337    -31
89338    -35
89339    -35
89340    -35
89341    -35
Name: hr, dtype: int64
```

Sorting

Pandas objects include methods for re-ordering data.

```
In [92]: baseball_newind.sort_index().head()

Out[92]:
```

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb |
|-------------------------|-----------|------|-------|------|----|-----|-----|----|-----|-----|-----|------|-----|-----|----|
| alomasa02NYN2007 | alomasa02 | 2007 | 1 | NYN | NL | 8 | 22 | 1 | 3 | 1 | ... | 0.0 | 0.0 | 0.0 | 0 |
| aloumo01NYN2007 | aloumo01 | 2007 | 1 | NYN | NL | 87 | 328 | 51 | 112 | 19 | ... | 49.0 | 3.0 | 0.0 | 27 |
| ausmubr01HOU2007 | ausmubr01 | 2007 | 1 | HOU | NL | 117 | 349 | 38 | 82 | 16 | ... | 25.0 | 6.0 | 1.0 | 37 |
| benitar01FLO2007 | benitar01 | 2007 | 2 | FLO | NL | 34 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 |
| benitar01SFN2007 | benitar01 | 2007 | 1 | SFN | NL | 19 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 |

5 rows × 22 columns

```
In [93]: baseball_newind.sort_index(ascending=False).head()

Out[93]:
```

| | player | year | stint | team | lg | g | ab | r | h | X2b | ... | rbi | sb | cs | bb |
|-------------------------|----------|------|-------|------|----|-----|-----|----|----|-----|-----|------|-----|-----|----|
| zaungr01TOR2007 | zaungr01 | 2007 | 1 | TOR | AL | 110 | 331 | 43 | 80 | 24 | ... | 52.0 | 0.0 | 0.0 | 51 |
| womacto01CHN2006 | | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | ... | 5.0 | 5.0 | 5.0 | 5 |

| | | | | | | | | | | | | | | | | |
|-------------------------|-----------|------|---|-----|----|----|----|---|---|---|-----|-----|-----|-----|---|---|
| witasja01TBA2007 | witasja01 | 2007 | 1 | TBA | AL | 3 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 0 |
| williwo02HOU2007 | williwo02 | 2007 | 1 | HOU | NL | 33 | 59 | 3 | 6 | 0 | ... | 2.0 | 0.0 | 0.0 | 0 | 2 |
| wickmbo01ATL2007 | wickmbo01 | 2007 | 1 | ATL | NL | 47 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0 | 0 |

5 rows × 22 columns

Try sorting the **columns** instead of the rows, in ascending order:

```
In [94]: baseball_newind.sort_index(axis=1).head()
```

```
Out[94]:
```

| | X2b | X3b | ab | bb | cs | g | gidp | h | hbp | hr | ... | player | r | rbi | sb | sf | sh | so |
|-------------------------|-----|-----|----|----|-----|---|------|---|-----|----|-----|--------|---|-----|-----|-----|-----|-----|
| womacto01CHN2006 | 5 | 5 | 5 | 5 | 5.0 | 5 | 5.0 | 5 | 5.0 | 5 | ... | 5 | 5 | 5.0 | 5.0 | 5.0 | 5.0 | 5.0 |
| schilcu01BOS2006 | 5 | 5 | 5 | 5 | 5.0 | 5 | 5.0 | 5 | 5.0 | 5 | ... | 5 | 5 | 5.0 | 5.0 | 5.0 | 5.0 | 5.0 |
| myersmi01NYA2006 | 5 | 5 | 5 | 5 | 5.0 | 5 | 5.0 | 5 | 5.0 | 5 | ... | 5 | 5 | 5.0 | 5.0 | 5.0 | 5.0 | 5.0 |
| helliri01MIL2006 | 5 | 5 | 5 | 5 | 5.0 | 5 | 5.0 | 5 | 5.0 | 5 | ... | 5 | 5 | 5.0 | 5.0 | 5.0 | 5.0 | 5.0 |
| johnsra05NYA2006 | 5 | 5 | 5 | 5 | 5.0 | 5 | 5.0 | 5 | 5.0 | 5 | ... | 5 | 5 | 5.0 | 5.0 | 5.0 | 5.0 | 5.0 |

5 rows × 22 columns

We can also use `sort_values` to sort a `Series` by value, rather than by label.

```
In [95]: baseball.hr.sort_values(ascending=False).head(10)
```

```
Out[95]:
```

| | |
|-------|----|
| id | |
| 89360 | 35 |
| 89462 | 30 |
| 89521 | 28 |
| 89361 | 26 |
| 89378 | 25 |
| 89489 | 24 |
| 89374 | 21 |
| 89371 | 21 |
| 89396 | 20 |
| 89439 | 20 |

Name: hr, dtype: int64

For a `DataFrame`, we can sort according to the values of one or more columns using the `by` argument of `sort_values`:

```
In [96]: baseball[['player', 'sb', 'cs']].sort_values(ascending=[False, True],
                                                    by=['sb', 'cs']).head(10)
```

```
Out[96]:
```

| | player | sb | cs |
|--------------|-----------|------|-----|
| id | | | |
| 89378 | sheffga01 | 22.0 | 5.0 |
| 89430 | loftoke01 | 21.0 | 4.0 |
| 89347 | vizquom01 | 14.0 | 6.0 |
| 89463 | greensh01 | 11.0 | 1.0 |
| 88652 | finlest01 | 7.0 | 0.0 |
| 89462 | griffke02 | 6.0 | 1.0 |
| 89530 | ausmubr01 | 6.0 | 1.0 |

| | | | |
|--------------|-----------|-----|-----|
| 89466 | gonzalu01 | 6.0 | 2.0 |
| 89521 | bondsba01 | 5.0 | 0.0 |
| 89438 | kleskry01 | 5.0 | 1.0 |

Exercise 5

Calculate **on base percentage** for each player, and return the ordered series of estimates.

$$obp = \frac{h + bb + hbp}{ab + bb + hbp + sf}$$

```
In [97]: baseball['obp']=baseball.apply(lambda p: (p.h+p.bb+p.hbp)/(p.ab+p.bb+p.hbp+p.sf) if (p.a
```

Missing data

The occurrence of missing data is so prevalent that it pays to use tools like Pandas, which seamlessly integrates missing data handling so that it can be dealt with easily, and in the manner required by the analysis at hand.

Missing data are represented in `Series` and `DataFrame` objects by the `NaN` floating point value. However, `None` is also treated as missing, since it is commonly used as such in other contexts (e.g. NumPy).

```
In [98]: foo = pd.Series([np.nan, -3, None, 'foobar'])
foo
```

```
Out[98]: 0      NaN
1       -3
2      None
3    foobar
dtype: object
```

```
In [99]: foo.isnull()
```

```
Out[99]: 0      True
1     False
2      True
3     False
dtype: bool
```

Missing values may be dropped or indexed out:

```
In [100... foo.dropna()
```

```
Out[100]: 1       -3
3    foobar
dtype: object
```

```
In [101... foo[foo.notnull()]
```

```
Out[101]: 1       -3
3    foobar
dtype: object
```

By default, `dropna` drops entire rows in which one or more values are missing.

```
In [102... data.dropna()
```

```
Out[102]:
```

| | patient | phylum | value | year | treatment |
|---|---------|----------------|-------|------|-----------|
| 0 | 1 | Firmicutes | 632 | 2013 | 0.0 |
| 1 | 1 | Proteobacteria | 1638 | 2013 | 0.0 |
| 2 | 1 | Actinobacteria | 569 | 2013 | 0.0 |
| 3 | 1 | Bacteroidetes | 14 | 2013 | 0.0 |
| 4 | 2 | Firmicutes | 21 | 2013 | 1.0 |
| 5 | 2 | Proteobacteria | 0 | 2013 | 1.0 |

This can be overridden by passing the `how='all'` argument, which only drops a row when every field is a missing value.

```
In [103... data.dropna(how='all')
```

```
Out[103]:
```

| | patient | phylum | value | year | treatment |
|---|---------|----------------|-------|------|-----------|
| 0 | 1 | Firmicutes | 632 | 2013 | 0.0 |
| 1 | 1 | Proteobacteria | 1638 | 2013 | 0.0 |
| 2 | 1 | Actinobacteria | 569 | 2013 | 0.0 |
| 3 | 1 | Bacteroidetes | 14 | 2013 | 0.0 |
| 4 | 2 | Firmicutes | 21 | 2013 | 1.0 |
| 5 | 2 | Proteobacteria | 0 | 2013 | 1.0 |
| 6 | 2 | Actinobacteria | 5 | 2013 | NaN |
| 7 | 2 | Bacteroidetes | 555 | 2013 | NaN |

Rather than omitting missing data from an analysis, in some cases it may be suitable to fill the missing value in, either with a default value (such as zero) or a value that is either imputed or carried forward/backward from similar data points. We can do this programmatically in Pandas with the `fillna` argument.

```
In [104... bacteria2.fillna(0)
```

```
Out[104]:
```

```
phylum
Firmicutes      0.0
Proteobacteria  632.0
Actinobacteria  1638.0
Bacteroidetes   569.0
dtype: float64
```

```
In [105... data.fillna({'year': 2013, 'treatment': 2})
```

```
Out[105]:
```

| | patient | phylum | value | year | treatment |
|---|---------|----------------|-------|------|-----------|
| 0 | 1 | Firmicutes | 632 | 2013 | 0.0 |
| 1 | 1 | Proteobacteria | 1638 | 2013 | 0.0 |
| 2 | 1 | Actinobacteria | 569 | 2013 | 0.0 |
| 3 | 1 | Bacteroidetes | 14 | 2013 | 0.0 |
| 4 | 2 | Firmicutes | 21 | 2013 | 1.0 |

| | | | | | |
|---|---|----------------|-----|------|-----|
| 5 | 2 | Proteobacteria | 0 | 2013 | 1.0 |
| 6 | 2 | Actinobacteria | 5 | 2013 | 2.0 |
| 7 | 2 | Bacteroidetes | 555 | 2013 | 2.0 |

Notice that `fillna` by default returns a new object with the desired filling behavior, rather than changing the `Series` or `DataFrame` **in place**.

We can alter values in-place using `inplace=True` .

```
In [106... data['treatment'].fillna(2, inplace=True)
data
```

```
Out[106]:
```

| | patient | phylum | value | year | treatment |
|---|---------|----------------|-------|------|-----------|
| 0 | 1 | Firmicutes | 632 | 2013 | 0.0 |
| 1 | 1 | Proteobacteria | 1638 | 2013 | 0.0 |
| 2 | 1 | Actinobacteria | 569 | 2013 | 0.0 |
| 3 | 1 | Bacteroidetes | 14 | 2013 | 0.0 |
| 4 | 2 | Firmicutes | 21 | 2013 | 1.0 |
| 5 | 2 | Proteobacteria | 0 | 2013 | 1.0 |
| 6 | 2 | Actinobacteria | 5 | 2013 | 2.0 |
| 7 | 2 | Bacteroidetes | 555 | 2013 | 2.0 |

Missing values can also be interpolated, using any one of a variety of methods:

```
In [107... bacteria2.fillna(method='bfill')
```

```
Out[107]:
```

| | |
|----------------|--------|
| phylum | |
| Firmicutes | 632.0 |
| Proteobacteria | 632.0 |
| Actinobacteria | 1638.0 |
| Bacteroidetes | 569.0 |

dtype: float64

Data summarization

We often wish to summarize data in `Series` or `DataFrame` objects, so that they can more easily be understood or compared with similar data. The NumPy package contains several functions that are useful here, but several summarization or reduction methods are built into Pandas data structures.

```
In [108... baseball.sum()
```

```
Out[108]:
```

| | |
|--------|---|
| player | womacto01schilcu01myersmi01helliri01johnsra05f... |
| year | 200692 |
| stint | 113 |
| team | CHNBOSNYAMILNYASFNARILANATLNYNTORTBAHOUARIATLM... |
| lg | NLALALNLALNLNLNLNLNLALALNLNLNLALNLNLNLNLALALNL... |
| g | 5238 |
| ab | 13654 |
| r | 1869 |
| h | 3582 |
| x2b | 739 |
| x3b | 55 |

```

hr          437
rbi         1847.0
sb          138.0
cs          46.0
bb          1549
so          2408.0
ibb         177.0
hbp         112.0
sh          138.0
sf          120.0
gidp        354.0
obp         20.751829
dtype: object

```

Clearly, `sum` is more meaningful for some columns than others. For methods like `mean` for which application to string variables is not just meaningless, but impossible, these columns are **automatically excluded**:

```
In [109... baseball.mean(numeric_only=True)
```

```

Out[109]: year      2006.920000
          stint     1.130000
          g         52.380000
          ab       136.540000
          r         18.690000
          h         35.820000
          X2b        7.390000
          X3b        0.550000
          hr         4.370000
          rbi       18.470000
          sb         1.380000
          cs         0.460000
          bb        15.490000
          so        24.080000
          ibb        1.770000
          hbp        1.120000
          sh         1.380000
          sf         1.200000
          gidp       3.540000
          obp        0.207518
          dtype: float64

```

The important difference between NumPy's functions and Pandas' methods is that the latter have built-in support for handling missing data.

```
In [110... bacteria2
```

```

Out[110]: phylum
          Firmicutes      NaN
          Proteobacteria  632.0
          Actinobacteria  1638.0
          Bacteroidetes   569.0
          dtype: float64

```

```
In [111... bacteria2.mean()
```

```
Out[111]: 946.3333333333334
```

Sometimes we may not want to ignore missing values, and allow the `nan` to propagate.

```
In [112... bacteria2.mean(skipna=False)
```

```
Out[112]: nan
```

Passing `axis=1` will summarize over rows instead of columns, which only makes sense in certain situations.

```
In [113]: extra_bases = baseball[['X2b', 'X3b', 'hr']].sum(axis=1)
extra_bases.sort_values(ascending=False).head(20)
```

```
Out[113]: id
88653      69
89439      57
89361      56
89462      55
89396      54
89360      54
89489      54
89371      50
89378      46
89374      46
89389      45
89523      44
89521      42
89463      41
89466      40
88652      39
89438      36
89330      35
89533      33
89481      29
dtype: int64
```

A useful summarization that gives a quick snapshot of multiple statistics for a `Series` or `DataFrame` is `describe`:

```
In [114]: baseball.describe()
```

```
Out[114]:
```

| | year | stint | g | ab | r | h | X2b | X3b |
|-------|------------|------------|------------|------------|-----------|------------|------------|------------|
| count | 100.00000 | 100.000000 | 100.000000 | 100.000000 | 100.00000 | 100.000000 | 100.000000 | 100.000000 |
| mean | 2006.92000 | 1.130000 | 52.380000 | 136.540000 | 18.69000 | 35.820000 | 7.390000 | 0.550000 |
| std | 0.27266 | 0.337998 | 48.031299 | 181.936853 | 27.77496 | 50.221807 | 11.117277 | 1.445124 |
| min | 2006.00000 | 1.000000 | 1.000000 | 0.000000 | 0.00000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 2007.00000 | 1.000000 | 9.500000 | 2.000000 | 0.00000 | 0.000000 | 0.000000 | 0.000000 |
| 50% | 2007.00000 | 1.000000 | 33.000000 | 40.500000 | 2.00000 | 8.000000 | 1.000000 | 0.000000 |
| 75% | 2007.00000 | 1.000000 | 83.250000 | 243.750000 | 33.25000 | 62.750000 | 11.750000 | 1.000000 |
| max | 2007.00000 | 2.000000 | 155.000000 | 586.000000 | 107.00000 | 159.000000 | 52.000000 | 12.000000 |

`describe` can detect non-numeric data and sometimes yield useful information about it.

```
In [115]: baseball.player.describe()
```

```
Out[115]: count      100
unique       82
top    wellsda01
freq         2
Name: player, dtype: object
```

We can also calculate summary statistics *across* multiple columns, for example, correlation and covariance.

$$\text{cov}(x, y) = \sum_i (x_i - \bar{x})(y_i - \bar{y})$$

```
In [116... baseball.hr.cov(baseball.X2b)
```

```
Out[116]: 69.07646464646454
```

$$\text{corr}(x, y) = \frac{\text{cov}(x, y)}{(n-1)s_x s_y} = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_i (x_i - \bar{x})^2 \sum_i (y_i - \bar{y})^2}}$$

```
In [117... baseball.hr.corr(baseball.X2b)
```

```
Out[117]: 0.7790615182539742
```

```
In [118... baseball.ab.corr(baseball.h)
```


```
Out[118]: 0.9942174036272377
```

Plotting

Sometimes the correlation between two attributes is obvious when we plot them in the 2-d space. Also other kinds of plots (e.g., histograms) help us understand whether our data follows a particular law (e.g., a power law). Thus, it is crucial, not only to know how to select the appropriate slice of our dataset, but also how to properly visualize it.


```
In [119... baseball.plot.scatter(x='hr', y='X2b')
```

```
Out[119]: <Axes: xlabel='hr', ylabel='X2b'>
```




```
In [120]: baseball.plot.scatter(x='ab', y='h')
```


```
Out[120]: <Axes: xlabel='ab', ylabel='h'>
```



Do you understand which of the above has better correlation?


```
In [121]: baseball['ab'].hist()
```

```
Out[121]: <Axes: >
```



```
In [122]: baseball['h'].hist()
```

<Axes: >



What do you observe in the distributions of the two attributes? Is it safe to say that they correlate based on the last two plots? Why we need also the scatterplots?

Writing Data to Files

As well as being able to read several data input formats, Pandas can also export data to a variety of storage formats. We will bring your attention to just a couple of these.

```
In [123... mb.to_csv("mb.csv")
```

The `to_csv` method writes a `DataFrame` to a comma-separated values (csv) file. You can specify custom delimiters (via `sep` argument), how missing values are written (via `na_rep` argument), whether the index is written (via `index` argument), whether the header is included (via `header` argument), among other options.

An efficient way of storing data to disk is in **binary format**. Pandas supports this using Python's built-in pickle serialization.

```
In [124... baseball.to_pickle("baseball pickle")
```

The complement to `to_pickle` is the `read_pickle` function, which restores the pickle to a `DataFrame` or `Series` :

```
In [125]: pd.read_pickle("baseball pickle").head()
```

[illegible]

| | | | | | | | | | | | | | | | | | | | |
|--------------|-----------|------|---|-----|----|----|----|---|----|---|-----|-----|-----|---|-----|-----|-----|-----|-----|
| 88641 | womacto01 | 2006 | 2 | CHN | NL | 19 | 50 | 6 | 14 | 1 | ... | 1.0 | 1.0 | 4 | 4.0 | 0.0 | 0.0 | 3.0 | 0.0 |
| 88643 | schilcu01 | 2006 | 1 | BOS | AL | 31 | 2 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 88645 | myersmi01 | 2006 | 1 | NYA | AL | 62 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 88649 | helliri01 | 2006 | 1 | MIL | NL | 20 | 3 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0 | 2.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 88650 | johnsra05 | 2006 | 1 | NYA | AL | 33 | 6 | 0 | 1 | 0 | ... | 0.0 | 0.0 | 0 | 4.0 | 0.0 | 0.0 | 0.0 | 0.0 |

5 rows × 23 columns

As Wes warns in his book [1], it is recommended that binary storage of data via pickle only be used as a temporary storage format, in situations where speed is relevant. This is because there is no guarantee that the pickle format will not change with future versions of Python.

References

[1] [Python for Data Analysis](#) Wes McKinney